

# Hybrid Distributed Programming

Mohamad Ibrahim Cheikh

November 30, 2020

## Introduction

Parallel computing refers to simulations in which multiple computational resources are implemented to solve a problem at the same time. Parallelism is usually achieved by dividing the problem into many smaller ones, in which each one is solved separately and also at the same time. Different forms of parallel programming has been developed like bit-level, instruction-level, data, and task parallelism. Some widely used parallel programming paradigms are based on message passing, such as MPI [3], which is suitable both on distributed memory (DM) and distributed-shared memory (DSM) systems architectures. Other popular models are based over the parallel compiler directives, such as OpenMP [4] and HPF [5], where the OpenMP takes advantage of the shared memory parallelism, and the HPF exploits the data parallelism.

The current report will implement a hybrid MPI/Openmp to solve a problem on a high performance machine like Beocat.

## Experimental Configuration

The supercomputer on which we are going to conduct our performance analysis is Beocat, which is a high performance computer cluster at the Kansas State University [1]. The performance analysis will be performed on the Elves nodes of Beocat [2]. Table 1 gives a brief description of the nodes configuration.

| Elve Nodes                |                             |                             |                             |                             |
|---------------------------|-----------------------------|-----------------------------|-----------------------------|-----------------------------|
| Nodes                     | 1-56                        | 57-72,77                    | 73-76,78, 79                | 80-85                       |
| Processors                | 2x 8-Core Xeon E5-2690      | 2x 10-Core Xeon E5-2690 v2  | 2x 10-Core Xeon E5-2690 v2  | 2x 10-Core Xeon E5-2690v2   |
| Ram                       | 64GB                        | 96GB                        | 384GB                       | 64GB                        |
| Hard Drive                | 1x 250GB 7,200 RPM SATA     | 1x 250GB 7,200 RPM SATA     | 1x 250GB 7,200 RPM SATA     | 1x 250GB 7,200 RPM SATA     |
| NICs                      | 4x Intel I350               | 4x Intel I350               | 4x Intel I350               | 4x Intel I350               |
| 10GbE and QDR Infini-band | MT27500 Family (ConnectX-3) | MT27500 Family (ConnectX-3) | MT27500 Family (ConnectX-3) | MT27500 Family (ConnectX-3) |

Table 1: Elves node configuration taken from [2]

## Part 1: Numerical Integration of Easom's function

The first problem that requires the implementation of the MPI/Openmp algorithm is finding the numerical integration of Easom's function defined below:

$$f(x, y) = -\cos(x)\sin(y)\exp(-[(x - \pi)^2 + (y - \pi)^2]) \quad (1)$$

To integrate this function one has first to discretize the domain into smaller sub domains. Since this is a 2D problem then the smaller domains will be rectangles of size  $dx.dy$ . After dividing the domain, the program has to loop over all the small rectangles, multiply them by  $f(x, y)$  and sum them up to get the the final integral value. So the code should do something like:

```
1 for(i = 0; i < endXdir ; i++)
2   for(j = 0; j < endYdir; j++)
3     volume += h*h*f(x + i*h, y + j*h);
```

where it loops over the points in  $x$  and  $y$ , calculates the volume using  $h$  (which is a small number  $\approx 0.00001$ ), and then add them up to get the final integral value.

### MPI/Openmp

The hybrid MPI/Openmp implementation for the current problem will try to divide the domain over all the cores. The first step it will do is divide the domain along the  $x$  direction according to the number of MPI processes. This is achieved by controlling the *starting* and *ending* of  $i$ -th loop or first loop. After the domain is divided into multiple smaller domains along the  $x$ -axis the code will enter the Openmp environment and discretize the mesh into even smaller domains, but this time along the  $y$  direction according to the number of Openmp threads.

So as a summary, we started with a large domain, and divided it first into smaller domains along the  $x$ -axis for the MPI environment, and then divided into smaller domains along the  $y$ -axis for the OpenMP domain. After the program finishes with summing up the volumes in the parallelized OpenMP environment it will add them up together to get the final OpenMP version, by utilizing the critical command in Openmp. When this is over each MPI process will have part of the whole volume, to get the final one we will use MPI-reduce to get the final value of the integration. The figure gives a better picture of how the code works.

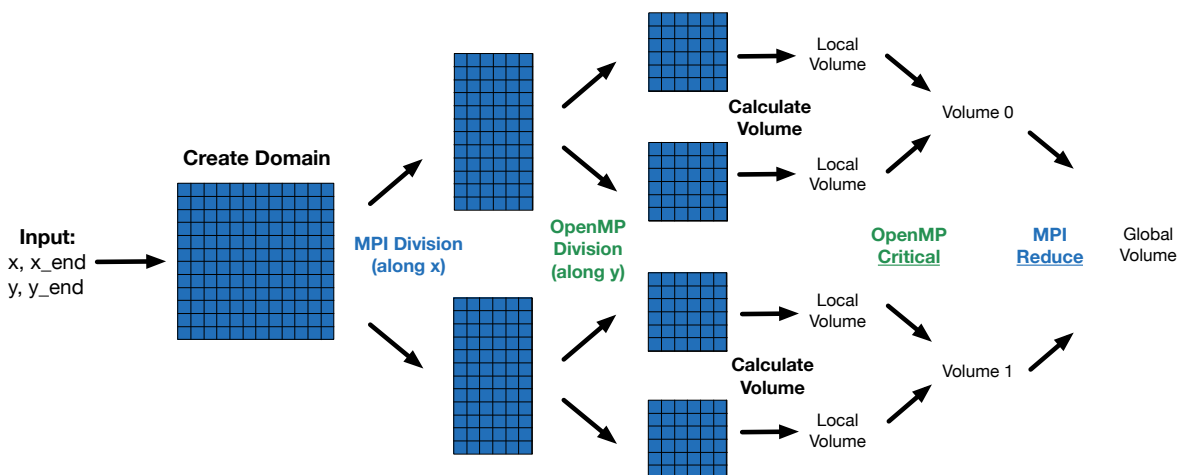


Figure 1: Hybrid MPI/Openmp implementation for Part 1, with MPI=2 and OpenMP=2

## Part 2: Monte Carlo Method

The second problem that requires the implementation of the hybrid MPI/Openmp algorithm is optimization using the Monte-Carlo Method. The function that will be analyzed is defined below:

$$f(x) = \cos(x) + |7.0 - x|^{2/15} + 2|5.0 - x|^{4/35} \quad (2)$$

The program will start by generate a set of random numbers between the first and last x values. After which the program will loop over the set of x values and calculate for each one the corresponding  $f(x)$ , and save the lowest one

### MPI/Openmp

The hybrid MPI/Openmp implementation for this problem will try to divide the number of points over all the cores. The first thing that the program will do is divide the total number of points over the MPI-Process. So if we requested that the program use 1000 points and we have 4 MPI process then each process will generate 250 random point. After generating the set of random points the program will enter the Openmp enviroment and start the Monte-Carlo Method. Inside the Monte-Carlo method the code will enter the Openmp environment and will divide the set into an even smaller set according to the number of Openmp threads, and find the value for each point. At the end of the Openmp environment the program will save the minimum value. After exiting the Openmp enviroment, the cores will send their minimum value to the MPI-Process of rank 0 so it can calculate the global minimum.

So as a summary, we started with a large set of points, and divid it first into smaller sets for each MPI process, and then divided more into smaller set for each OpenMp thread. The figure gives a better picture of how the code works.

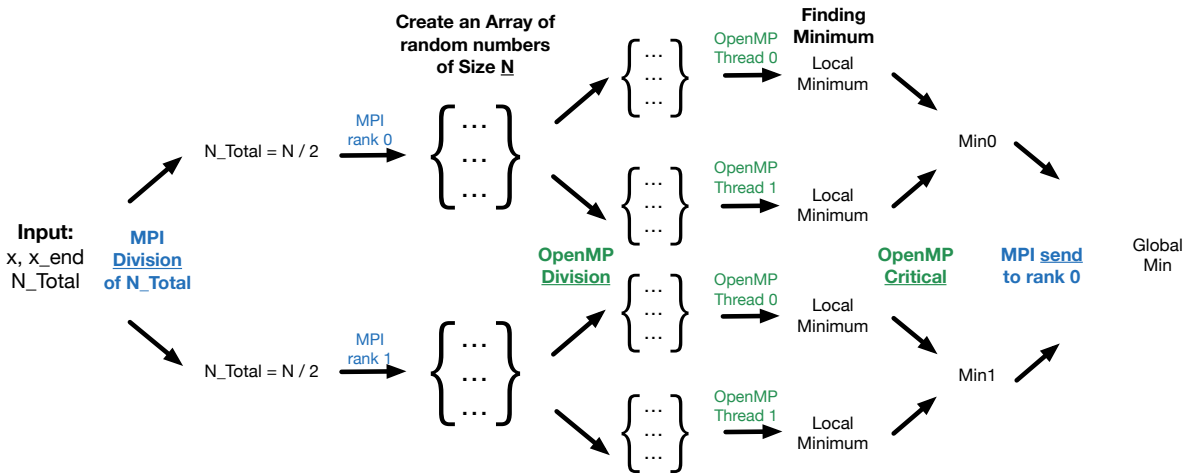


Figure 2: Hybrid MPI/Openmp implementation for Part 2, with MPI=2 and OpenMP=2

## Part 3: Compute the standard deviation

The third problem that requires the implementation of the MPI/Openmp algorithm is a program that computes the standard deviation of a set of randomly generated numbers. The equation that governs the standard deviation  $\sigma$  of a probability distribution is defined as the square root of the variance  $\sigma^2$ ,

$$\sigma = \sqrt{\langle x^2 \rangle + \langle x \rangle^2} \quad (3)$$

where  $\langle x^2 \rangle$  sums of squares and  $\langle x \rangle$  is the sum of all the numbers.

## MPI/Openmp

The hybrid MPI/Openmp implementation for this problem is similar to the previous problem (Part 2). The program will first divide the total number of points over the MPI-Process. So if we requested that the program use 1000 points and we have 4 MPI process then each process will generate 250 random point. After generating the set of random points the program will enter the Openmp enviroment and start the *Standard Deviation* function. Inside the function the code will enter the Openmp environment and will divide the set into an even smaller set according to the number of Openmp threads. Each Openmp thread will calculate its local value  $\sum x^2$  and  $\sum x$ , and at the end the Openmp will add all the local values together. When the calculation ends, the cores will send their value of  $\sum x^2$  and  $\sum x$  to rank 0 so it can add them up and find their mean, as follows

$$\langle x \rangle = \frac{\sum_{i=0}^{n=cores} (\sum x)_i}{\text{Total Num of Points}} \quad \text{and} \quad \langle x^2 \rangle = \frac{\sum_{i=0}^{n=cores} (\sum x^2)_i}{\text{Total Num of Points}} \quad (4)$$

finally rank 0 will calculate the standard deviation as follows

$$\sigma = \sqrt{\langle x^2 \rangle + \langle x \rangle^2} \quad (5)$$

Figure gives a better picture of how the code works.

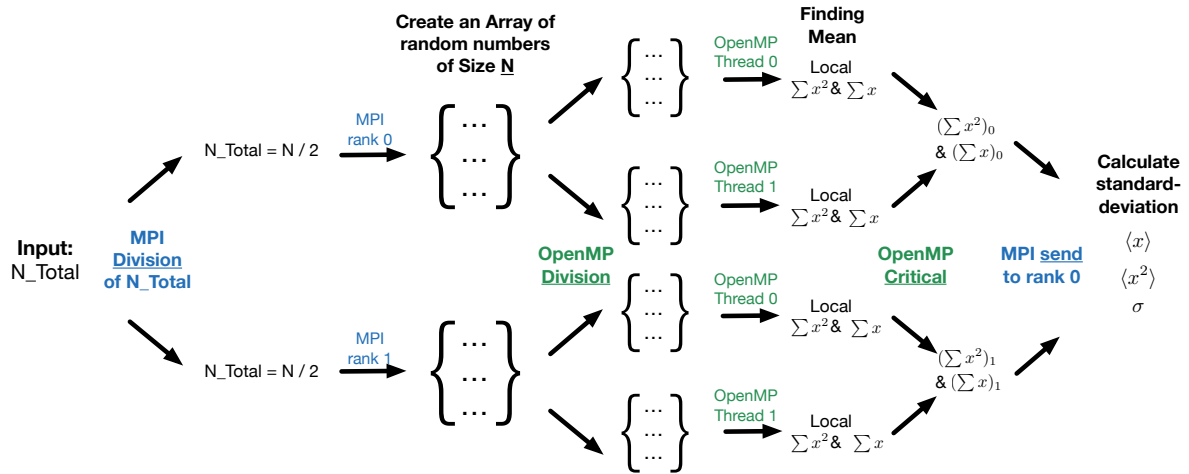


Figure 3: Hybrid MPI/Openmp implementation for Part 3, with MPI=2 and OpenMP=2

## Compiling the three programs

The three programs were compiled on beocat using the command:

```
1 mpicc -O2 -fopenmp -o Part1-MPI-OpenMP Part1-MPI-OpenMP.c -lm
```

where a level 2 optimization was implemented.

## Running the Hybrid MPI/OpenMP Programs on Beocat

The implementation of the Hybrid MPI/OpenMP was proved to be difficult. For the case of OpenMP=1 and MPI=N, the implementation was easy, all we had to do is write the following line in the bash script

```
1 $$ -pe mpi-spread N
2 mpirun -np N ./program 1 ## 1 for the OpenMP threads
```

An example of how the output should be for Part 1 with MPI=16, and OpenMP=1 is shown below

---

```
1 Thread 0 out of 1 from process 0 out of 16 on elf86
2 Thread 0 out of 1 from process 1 out of 16 on elf93
3 Thread 0 out of 1 from process 15 out of 16 on elf79
4 Thread 0 out of 1 from process 2 out of 16 on elf92
5 Thread 0 out of 1 from process 11 out of 16 on elf90
6 Thread 0 out of 1 from process 10 out of 16 on elf91
7 Thread 0 out of 1 from process 5 out of 16 on elf96
8 Thread 0 out of 1 from process 12 out of 16 on elf89
9 Thread 0 out of 1 from process 7 out of 16 on elf94
10 Thread 0 out of 1 from process 9 out of 16 on elf99
11 Thread 0 out of 1 from process 14 out of 16 on elf87
12 Thread 0 out of 1 from process 13 out of 16 on elf88
13 Thread 0 out of 1 from process 3 out of 16 on elf03
14 Thread 0 out of 1 from process 8 out of 16 on elf05
15 Thread 0 out of 1 from process 6 out of 16 on elf95
16 Thread 0 out of 1 from process 4 out of 16 on elf98
```

---

However for the case with OpenMP > 1 it is much more difficult. For example let us assume a case with MPI=4 and OpenMP = 4. The first thing we have to do is assign the correct number of cores, this is done by writing the command

---

```
1 # $ -pe mpi-4 16
```

---

This will assign 4 machines, with 4 cores on each machine. The next step is we have to update the host file and make it so that the number of slots on each machine is 1. To do that we added the lines below into the bash file that make use of the command *sed*:

---

```
1 sed 's/ 4/ slots=1/g' $PE_HOSTFILE > NEW_HOSTFILE3
2 sed -i 's/gen-reserved.q@elf...beocat.ksu.edu UNDEFINED//g'
   NEW_HOSTFILE3
```

---

The first line will search for the number "4" and replace it with "slots=1", and save it in a **new hostfile** called "NEW\_HOSTFILE3". As for the second line it will clean the data after the "slots=1" command. So the **new hostfile** will look like

---

```
1 elf93.beocat.ksu.edu slots=1
2 elf92.beocat.ksu.edu slots=1
3 elf03.beocat.ksu.edu slots=1
4 elf98.beocat.ksu.edu slots=1
```

---

while the **old hostfile** looked like

---

```
1 elf93.beocat.ksu.edu 4 gen-reserved.q@elf93.beocat.ksu.edu UNDEFINED
2 elf92.beocat.ksu.edu 4 gen-reserved.q@elf92.beocat.ksu.edu UNDEFINED
3 elf03.beocat.ksu.edu 4 gen-reserved.q@elf...beocat.ksu.edu UNDEFINED
4 elf98.beocat.ksu.edu 4 gen-reserved.q@elf...beocat.ksu.edu UNDEFINED
```

---

After creating the **new hostfile** we will give it as input to the mpi command by typing

---

```
1 mpirun --hostfile NEW_HOSTFILE3 -np 4 ./program 4 ## 4 = num OpenMP
```

---

An example of how the output should be for Part 1 with MPI=4, and OpenMP=4 is shown below

---

```
1 Thread 2 out of 4 from process 0 out of 4 on elf93
2 Thread 3 out of 4 from process 0 out of 4 on elf93
3 Thread 0 out of 4 from process 0 out of 4 on elf93
4 Thread 1 out of 4 from process 0 out of 4 on elf93
```

---

```

5 Thread 0 out of 4 from process 1 out of 4 on elf92
6 Thread 3 out of 4 from process 1 out of 4 on elf92
7 Thread 1 out of 4 from process 1 out of 4 on elf92
8 Thread 2 out of 4 from process 1 out of 4 on elf92
9 Thread 0 out of 4 from process 3 out of 4 on elf98
10 Thread 1 out of 4 from process 3 out of 4 on elf98
11 Thread 2 out of 4 from process 3 out of 4 on elf98
12 Thread 3 out of 4 from process 3 out of 4 on elf98
13 Thread 2 out of 4 from process 2 out of 4 on elf03
14 Thread 3 out of 4 from process 2 out of 4 on elf03
15 Thread 1 out of 4 from process 2 out of 4 on elf03
16 Thread 0 out of 4 from process 2 out of 4 on elf03

```

---

For a look at the full bash script please refer to the appendix

## Results

The three problems discussed in the previous section were tested for different numbers of inputs and with different number of cores (1, 2, 4, 8, and 16). The table below summarizes the number of cores used: The maximum OpenMP thread was 8.

| Summary of the Cores Used |             |                |
|---------------------------|-------------|----------------|
| Total Nodes               | MPI Threads | OpenMP Threads |
| 16                        | 16          | 1              |
| 16                        | 8           | 2              |
| 16                        | 4           | 4              |
| 16                        | 2           | 8              |
| 8                         | 8           | 1              |
| 8                         | 4           | 2              |
| 8                         | 2           | 4              |
| 8                         | 1           | 8              |
| 4                         | 4           | 1              |
| 4                         | 2           | 2              |
| 4                         | 1           | 4              |
| 2                         | 2           | 1              |
| 2                         | 1           | 2              |
| 1                         | 1           | 1              |

The results of part 1, 2, and 3 are shown in figures 4, 5 and 6.

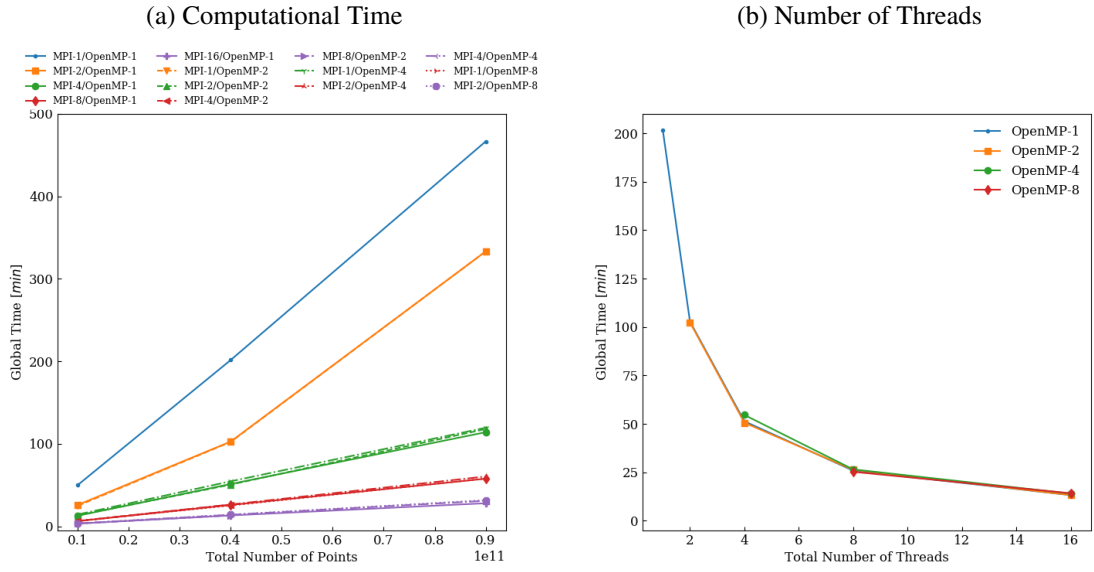


Figure 4: Results of Part 1; (a) Computational time versus number of input, (b) Computational Time versus total number of threads for a single input




Figure 5: Results of Part 2; (a) Computational time versus number of input, (b) Computational Time versus total number of threads for a single input

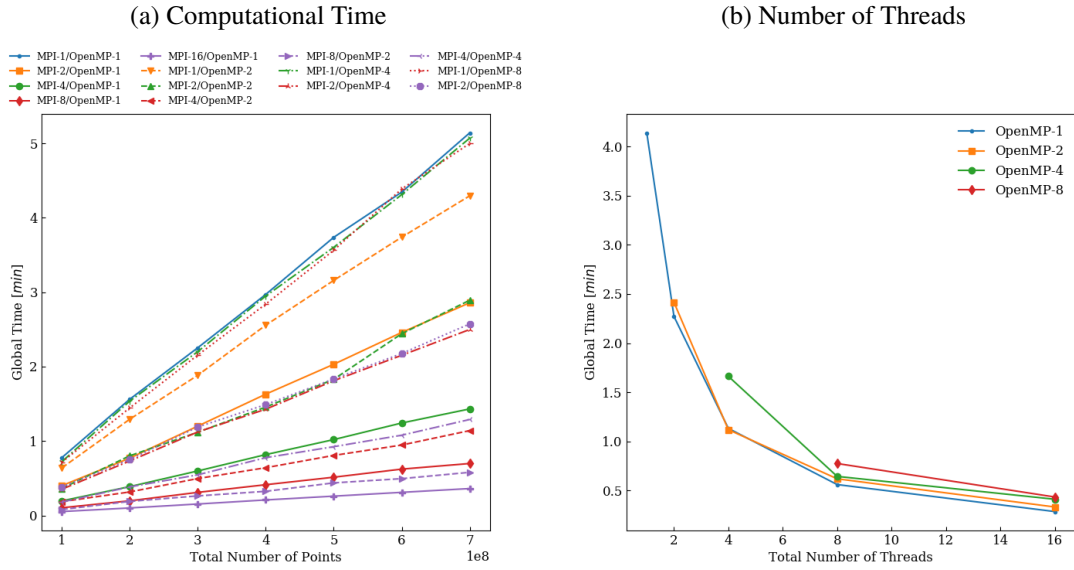


Figure 6: Results of Part 3; (a) Computational time versus number of input, (b) Computational Time versus total number of threads for a single input

From all the results it is clear that the pure MPI was faster than the hybrid OpenMP. At first this was surprising, since OpenMP utilizes the same *Shared memory* which is usually more efficient than the MPI *message passing*, as the latter usually requires increased data movement (moving data from the source to its destination) which is costly both performance-wise and energy-wise.

So to analyze this problem we looked at all the problems, and all of them have one thing in common, they are **extremely parallel**. What I mean is that the program can be easily parallelized without the need to create a difficult network between the threads. Thus the programs can easily divide the task according to the number of cores it has, and solve it without requiring the assistance of any of the other cores. The only time it requires communication is at the beginning, when the input are given, and at the end to sum up the data.

Thus by using the OpenMP environment we are increasing the communication time, since the program has to create private variables, solve them, and then combine them together, using a messaging interface. While the pure MPI will solve the problem with less communication. Another drawback of OpenMP is the "Load balance", where a barrier is used to force the fastest cores to wait for the slowest cores in the OpenMP region, especially at the critical location.

The only disadvantage of the pure MPI environment is that it needs a larger memory requirement, then a hybrid MPI/Openmp.

## 1 References:

1. <https://support.beocat.ksu.edu/BeocatDocs/index.php/MainPage>
2. [https://support.beocat.ksu.edu/BeocatDocs/index.php/Compute\\_Nodes](https://support.beocat.ksu.edu/BeocatDocs/index.php/Compute_Nodes)
3. W. D. Gropp. Parallel computing and domain decomposition. In Fifth International Symposium on Domain Decomposition Methods for Partial Differential Equations, Philadelphia, PA, 1992.
4. B. Chapman, G. Jost, and R. Van Der Pas. Using OpenMP: portable shared memory parallel programming, volume 10. MIT press, 2008.
5. C. H. Koelbel, D. B. Loveman, R. S. Schreiber, G. L. Steele Jr, M. E. Zosel, D. E. Ulberg, A. J. Mallinckrodt, S. McKay, et al. The high performance fortran handbook. Computers in Physics, 8(4):428–428, 1994.



## 2 Appendix:

Listing 1: 'Part1-MPI-OpenMP.c'

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <math.h>
4  #include <sys/resource.h>
5  #include <mpi.h>
6  #include <time.h>
7  #include <omp.h>
8
9  #define PI 3.14159265
10
11 double tstart, ttotal;
12 double volume ,global_volume;
13 int MPI_nthreads = 1, omp_nthreads;
14 double x,y,x_end,y_end;
15 char processor_name[20];
16
17 /* myclock: (Calculates the time)
18 -----*/
19 double myclock()
20 {
21     static time_t t_start = 0;  // Save and subtract off each time
22
23     struct timespec ts;
24     clock_gettime(CLOCK_REALTIME, &ts);
25     if( t_start == 0 ) t_start = ts.tv_sec;
26
27     return (double) (ts.tv_sec - t_start) + ts.tv_nsec * 1.0e-9;
28 }
29
30
31 /* Easom's Function
32 -----*/
33 double f (double x, double y)
34 {
35     return -cos(x)*sin(y)*exp( -((x - PI)*(x - PI) + (y - PI)*(y - PI)) )
36     ;
37 }
38
39 /* Calculate the Integral of the list
40 -----*/
41
42 void * Integrate (void *rank_MPI, void* rank_OpenMP)
43 {
44
45     double local_volume;
46     double h = 0.00001;
47
48     // MPI and OpenMP ID's
49     int myMPI_ID = *((int*) rank_MPI);
50     int myOp_ID;
51 }
```

```

52 // Total number of points
53 long N_i = (x_end - x) / h;
54 long N_j = (y_end - y) / h;
55
56 // Start i and End i controlled by MPI
57 int i_start = (myMPI_ID) * (N_i / MPI_nthreads);
58 int i_end = i_start + (N_i / MPI_nthreads);
59
60 // Start j and End j controlled by OpenMp
61 int j_start, j_end;
62
63 int i,j;
64
65 #pragma omp private(rank_OpenMP,j_start,j_end,local_volume,i,j)
66 {
67
68     printf("Thread %d out of %d from process %d out of %d on %s\n",
69         omp_get_thread_num(), omp_get_num_threads(), myMPI_ID,
70         MPI_nthreads, processor_name);
71
72     myOp_ID = ((int) rank_OpenMP);
73
74     local_volume = 0.0;
75
76     // Start i and End i controlled by MPI
77     j_start = ( myOp_ID) * (N_j / omp_nthreads);
78     j_end = j_start + (N_j / omp_nthreads);
79
80     // Initialize Local Integration
81     for(i = i_start; i < i_end; i++)
82     {
83         for(j = j_start; j < j_end; j++)
84         {
85             local_volume += h*h*f(x + i*h + h/2.0, y + j*h + h/2.0);
86         }
87     }
88
89     #pragma omp critical
90     volume += local_volume;
91 }
92 }
93
94 /* Main
95 -----*/
96 int main(int argc, char *argv[])
97 {
98     volume ,global_volume = 0.0, 0.0;
99     struct rusage ru;
100     // Default Value
101     omp_nthreads = 1;
102     x , y , x_end ,y_end = 1.0 ,1.0,5.0,5.0;
103     if (argc >= 2){
104         omp_nthreads = atoi(argv[1]);
105         sscanf(argv[2], "%lf",&x);

```

```

106     sscanf(argv[3], "%lf", &y);
107     sscanf(argv[4], "%lf", &x_end);
108     sscanf(argv[5], "%lf", &y_end);
109 }
110
111 // Initialize the MPI Enviroment
112 int i, rc, rank, namelen;
113
114 MPI_Status Status;
115 MPI_Request Request;
116 rc = MPI_Init(&argc, &argv);
117 if (rc != MPI_SUCCESS){
118     printf ("Error starting MPI program. Terminating.\n");
119     MPI_Abort(MPI_COMM_WORLD, rc);
120 }
121 MPI_Comm_size(MPI_COMM_WORLD, &MPI_nthreads); // Number of cores
122 MPI_Comm_rank(MPI_COMM_WORLD, &rank); // rank of each core
123 MPI_Get_processor_name(processor_name, &namelen); // Processors name
124
125 // Initialize the Open Mp Enviroment
126 omp_set_num_threads(omp_nthreads);
127
128 printf("rank:%d,x0:%lf,x_end:%lf,y0:%lf,y_end:%lf\n", rank, x, x_end, y,
        y_end);
129
130 if (rank == 0)
131 {
132     printf("Start MPI/OpenMP Process:\n");
133     tstart = myclock(); // Global Clock
134 }
135
136 // Calculate Integral
137 #pragma omp parallel
138 {
139     Integrate ( &rank, omp_get_thread_num() );
140 }
141
142
143 getrusage(RUSAGE_SELF, &ru);
144 long MEMORY_USAGE = ru.ru_maxrss; // Memory usage in Kb
145
146 MPI_Reduce(&volume, &global_volume, 1, MPI_DOUBLE, MPI_SUM, 0,
        MPI_COMM_WORLD);
147
148 if(rank==0)
149 {
150     tttotal = myclock() - tstart;
151     long N_j = (y_end - y) / 0.00001;
152     printf("Part1-MPI-%d/OpenMP-%d:%ld,%lf,%lf,%lf,%lf,%.12lf,%lf,%ld\n",
        MPI_nthreads, omp_nthreads, N_j, x, x_end, y, y_end, global_volume,
        tttotal, MEMORY_USAGE);
153     printf("-----\n");
154 }
155 MPI_Finalize();
156

```

```

157     return 0;
158 }

```

---

Listing 2: 'Part2-MPI-OpenMP.c'

---

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <math.h>
4  #include <sys/resource.h>
5  #include <mpi.h>
6  #include <time.h>
7  #include <omp.h>
8
9  #define PI 3.14159265
10
11 #define MIN(a,b) (((a)<(b))?(a):(b))
12 #define MAX(a,b) (((a)>(b))?(a):(b))
13
14 double tstart, ttotal;
15 int MPI_nthreads = 1, omp_nthreads;
16 double x_0, x_end;
17 long N_Total, N; // Number of points
18 double min_f;
19
20 double *number_array;
21 /* myclock: (Calculates the time)
22 -----*/
23 double myclock()
24 {
25     static time_t t_start = 0; // Save and subtract off each time
26
27     struct timespec ts;
28     clock_gettime(CLOCK_REALTIME, &ts);
29     if( t_start == 0 ) t_start = ts.tv_sec;
30
31     return (double) (ts.tv_sec - t_start) + ts.tv_nsec * 1.0e-9;
32 }
33
34
35 /* Creates a list of Random numbers from x_0 to x_end
36 -----*/
37 void * Create_Number_Arrays(long n, int myID)
38 {
39     long i;
40     srand(time(NULL)+ (int) myID); // randomize seed
41     int Z = 0;
42     // allocate an array of size n pointers to chars
43     number_array = malloc(sizeof(double *)*n);
44     for ( i=0; i < n; i++)
45     {
46         Z = rand()% 1000000;
47         number_array[i] = x_0 + (Z/1000000.0)*(x_end - x_0); // random
            number from x_0 to X_end
48     }
49 }
50

```

```

51
52 /* Function that we wish to optimize
53 -----*/
54 double f (double x)
55 {
56     return cos(x)+(pow(fabs(7.0-x), 2.0/15.0))+2*(pow(fabs(5.0-x), 4.0/35
57         .0));
58 }
59
60 /* Monte Carlo Method
61 -----*/
62
63 void * Monte_Carlo_Method (void* rank_OpenMP)
64 {
65
66     min_f = f(x_0);
67     double local_min_f;
68
69     // OpenMP ID's
70     int myOp_ID;
71
72     // Start j and End j controlled by OpenMp
73     int i, i_start, i_end;
74     double x;
75
76     #pragma omp private(myOp_ID,i_start,i_end,i,local_min_f,x)
77     {
78
79         myOp_ID = ((int) rank_OpenMP);
80         local_min_f = f(x_0);
81
82         // Start i and End i controlled by MPI
83         i_start = ( myOp_ID) * (N / omp_nthreads);
84         i_end = i_start + (N / omp_nthreads);
85
86         // Initialize Local Integration
87         for(i = i_start; i < i_end; i++)
88         {
89             x = number_array[i];
90             local_min_f = MIN(f(x),local_min_f);
91         }
92
93         #pragma omp critical
94         {
95             if (local_min_f < min_f)
96                 min_f = local_min_f;
97         }
98     }
99
100 }
101
102 /* Main
103 -----*/
104 int main(int argc, char *argv[])

```

```

105 {
106     struct rusage ru;
107     // Default Value
108     omp_nthreads = 1;
109     x_0 ,x_end = 0.0,10.0;
110     N_Total = 10;
111     if (argc >= 2){
112         omp_nthreads = atoi(argv[1]);
113         sscanf(argv[2], "%lf", &x_0);
114         sscanf(argv[3], "%lf", &x_end);
115         sscanf(argv[4], "%ld", &N_Total); // Number of points
116     }
117
118     // Initialize the MPI Enviroment
119     int i, rc, rank;
120
121     MPI_Status Status;
122     MPI_Request Request;
123     rc = MPI_Init(&argc,&argv);
124     if (rc != MPI_SUCCESS){
125         printf ("Error starting MPI program. Terminating.\n");
126         MPI_Abort(MPI_COMM_WORLD, rc);
127     }
128     MPI_Comm_size(MPI_COMM_WORLD,&MPI_nthreads); // Number of cores
129     MPI_Comm_rank(MPI_COMM_WORLD,&rank); // rank of each core
130
131     // Initialize the Open Mp Enviroment
132     omp_set_num_threads(omp_nthreads);
133
134     N = N_Total / MPI_nthreads; // number of points for each MPI
        Process
135
136     if (rank ==0)
137     {
138         printf("Start MPI/OpenMP Process:\n");
139         tstart = myclock(); // Global Clock
140     }
141
142     // Create the array of random numbers from x_0 to x_end
143     Create_Number_Arrays(N,rank);
144
145     // Find the Min using monte carlo for each MPI Process
146
147     #pragma omp parallel
148     {
149         Monte_Carlo_Method ( omp_get_thread_num() );
150     }
151
152     getrusage(RUSAGE_SELF, &ru);
153     long MEMORY_USAGE = ru.ru_maxrss; // Memory usage in Kb
154
155     // Send the data to MPI
156     if (rank != 0)
157     {
158         MPI_Isend(&min_f , 1, MPI_DOUBLE, 0, 1234, MPI_COMM_WORLD, &Request

```

```

    );
159     printf("Done rank %d\n",rank);
160 }
161
162 MPI_Barrier( MPI_COMM_WORLD ) ;
163
164 if (rank == 0)
165 {
166     double local_min_f=0.0;
167
168     for (i = 1; i< MPI_nthreads; i++)
169     {
170         MPI_Irecv(&local_min_f, 1, MPI_DOUBLE, i, 1234, MPI_COMM_WORLD, &
            Request);
171         min_f = MIN(local_min_f,min_f);
172     }
173
174     tttotal = myclock() - tstart;
175     printf("Part2-MPI-%d/OpenMP-%d:%ld,%lf,%lf,%ld\n",MPI_nthreads,
        omp_nthreads,N_Total,min_f,tttotal,MEMORY_USAGE);
176 }
177
178 MPI_Finalize();
179 return 0;
180 }

```

---

Listing 3: 'Part3-MPI-OpenMP.c'

---

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <math.h>
4  #include <sys/resource.h>
5  #include <mpi.h>
6  #include <time.h>
7  #include <omp.h>
8
9  double tstart, tttotal;
10 long N, N_total, sums, sums_squares, sums_recieved,
    sums_squares_recieved;
11 double stdev;
12 int MPI_nthreads = 1, omp_nthreads;
13
14 int *number_array; // A 1D array of char arrays (a pointer to pointers
    to chars)
15
16
17 /* myclock: (Calculates the time)
18 -----*/
19 double myclock()
20 {
21     static time_t t_start = 0; // Save and subtract off each time
22
23     struct timespec ts;
24     clock_gettime(CLOCK_REALTIME, &ts);
25     if( t_start == 0 ) t_start = ts.tv_sec;
26

```

```

27         return (double) (ts.tv_sec - t_start) + ts.tv_nsec * 1.0e-9;
28     }
29
30
31     /* Creates a list of Random numbers
32     -----*/
33     void * Create_Number_Arrays(int myID)
34     {
35         time_t t;
36         long i;
37         srand(time(NULL)+ (int) myID); // randomize seed
38         // allocate an array of size N pointers to chars
39         number_array = malloc(sizeof(int *)*N);
40
41         for ( i=0; i < N; i++)
42         {
43             number_array[i] = rand() % 1000; // random number from 0 to 1000
44         }
45         sums = 0;
46         sums_squares = 0;
47     }
48
49
50     // Print the list
51     void print_arrays()
52     {
53         long i;
54         for ( i=0; i < N; i++)
55         {
56             printf("%d\n",number_array[i] ); // random number from 0 to 1
57             million
58         }
59
60
61     /* Calculare the Sum of the list
62     -----*/
63     void * Calculate_Sums(void *myID)
64     {
65         long i;
66         long local_sums;
67         long local_sums_squares;
68         int startPos, endPos;
69
70         #pragma omp private(myID,startPos,endPos,local_sums,
71             local_sums_squares,i)
72         {
73             startPos = ((int) myID) * (N / omp_nthreads);
74             endPos = startPos + (N / omp_nthreads);
75
76             // Initialize Local sums
77             local_sums = 0;
78             local_sums_squares = 0;
79
79             for ( i= startPos; i < endPos; i++)

```



```

80     {
81         local_sums+= number_array[i] ; // random number from 0 to 1
            million
82         local_sums_squares += number_array[i]*number_array[i];
83     }
84
85     //printf("%d-%ld-%d-%d\n",omp_get_thread_num(),local_sums,startPos,
        endPos);
86
87     // sum up the partial sum into the global sum
88     #pragma omp critical
89     sums += local_sums;
90     sums_squares += local_sums_squares;
91 }
92 }
93
94 /* Calculate the Standard Deviation
95 -----*/
96 void * Calculate_Standard_Deviation()
97 {
98     stdev = sqrt( ((double)sums_squares/N_total) - ((double)sums/
        N_total)*((double)sums/ N_total) );
99 }
100
101 /* Main
102 -----*/
103 int main(int argc, char *argv[])
104 {
105
106     struct rusage ru;
107     // Default Value
108     N_total = 10;
109     omp_nthreads = 1;
110     if (argc >= 2){
111         omp_nthreads = atoi(argv[1]);
112         N_total = atoi(argv[2]);
113     }
114
115     // Initalize the MPI Enviroment
116     int i, rc, rank;
117
118     MPI_Status Status;
119     MPI_Request Request;
120     rc = MPI_Init(&argc,&argv);
121     if (rc != MPI_SUCCESS){
122         printf ("Error starting MPI program. Terminating.\n");
123         MPI_Abort(MPI_COMM_WORLD, rc);
124     }
125     MPI_Comm_size(MPI_COMM_WORLD,&MPI_nthreads); // Number of cores
126     MPI_Comm_rank(MPI_COMM_WORLD,&rank); // rank of each core
127
128     // Initalize the Open Mp Enviroment
129     omp_set_num_threads(omp_nthreads);
130
131     if (rank ==0)

```

```

132     {
133         printf("Start MPI/OpenMP Process:\n");
134         tstart = myclock(); // Global Clock
135     }
136     // Create Number Arrays
137     // -----
138
139     MPI_Bcast(&N_total, 1, MPI_LONG, 0, MPI_COMM_WORLD);
140
141     N = N_total / MPI_nthreads;
142
143     Create_Number_Arrays(rank); // All mpi process will create a random
        list
144     //print_arrays();
145     // Calculate Sums
146     #pragma omp parallel
147     {
148         Calculate_Sums(omp_get_thread_num());
149     }
150
151     getrusage(RUSAGE_SELF, &ru);
152     long MEMORY_USAGE = ru.ru_maxrss; // Memory usage in Kb
153
154
155     if (rank != 0)
156     {
157         MPI_Isend(&sums, 1, MPI_LONG, 0, 1234, MPI_COMM_WORLD, &Request);
158         MPI_Isend(&sums_squares, 1, MPI_LONG, 0, 5678, MPI_COMM_WORLD, &
            Request);
159         printf("Done rank %d\n",rank);
160     }
161     else
162     {
163
164         for (i = 1; i < MPI_nthreads; i++)
165         {
166             MPI_Irecv(&sums_recieved, 1, MPI_LONG, i, 1234, MPI_COMM_WORLD, &
                Request);
167             sums += sums_recieved ;
168
169             MPI_Irecv(&sums_squares_recieved, 1, MPI_LONG, i, 5678,
                MPI_COMM_WORLD, &Request);
170             sums_squares += sums_squares_recieved ;
171         }
172
173         Calculate_Standard_Deviation();
174         tttotal = myclock() - tstart;
175         printf("Part3-MPI-%d/OpenMP-%d:%ld,%lf,%lf,%ld\n",MPI_nthreads,
            omp_nthreads,N_total,stdev,tttotal,MEMORY_USAGE);
176     }
177
178     MPI_Finalize();
179     return 0;
180 }

```

---

Listing 4: Part 1 - Bash script for MPI equal 1

---

```

1  #!/bin/bash
2  $$ -l mem=5G
3  $$ -l h_rt=24:00:00
4  $$ -l killable
5  $$ -cwd
6  $$ -P KSU-GEN-RESERVED
7  $$ -q \*@@elves
8  $$ -pe single 8
9
10 # MPI = 1
11 for i in 2 3 4
12 do
13 mpirun -np 1 /homes/mcheikh/CIS_625/hw5/part-1/Part1-MPI-OpenMP 8 1 1
    $i $i
14 mpirun -np 1 /homes/mcheikh/CIS_625/hw5/part-1/Part1-MPI-OpenMP 1 1 1
    $i $i
15
16 echo -e "-----DONE-----\n"
17 done

```

---

Listing 5: Part 1 - Bash script for MPI equal 2

---

```

1  #!/bin/bash
2  $$ -l mem=5G
3  $$ -l h_rt=24:00:00
4  $$ -l killable
5  $$ -cwd
6  $$ -P KSU-GEN-RESERVED
7  $$ -q \*@@elves
8  $$ -pe mpi-8 16
9
10 echo "PE_HOSTFILE:"
11 echo $PE_HOSTFILE
12 echo
13 echo "cat PE_HOSTFILE:"
14 cat $PE_HOSTFILE
15 echo '-----'
16 sed 's/ 8/ slots=1/g' $PE_HOSTFILE > NEW_HOSTFILE0
17 sed -i 's/gen-reserved.q@elf...beocat.ksu.edu UNDEFINED//g'
    NEW_HOSTFILE0
18
19 # OpenMp = 2
20 for i in 2 3 4
21 do
22 mpirun --hostfile NEW_HOSTFILE0 -np 2 /homes/mcheikh/CIS_625/hw5/part
    -1/Part1-MPI-OpenMP 8 1 1 $i $i
23 mpirun --hostfile NEW_HOSTFILE0 -np 2 /homes/mcheikh/CIS_625/hw5/part
    -1/Part1-MPI-OpenMP 4 1 1 $i $i
24 mpirun --hostfile NEW_HOSTFILE0 -np 2 /homes/mcheikh/CIS_625/hw5/part
    -1/Part1-MPI-OpenMP 2 1 1 $i $i
25
26 echo -e "-----DONE-----\n"
27 done

```

---

Listing 6: Part 1 - Bash script for OpenMP equal 1

---

```

1  #!/bin/bash
2  $$ -l mem=5G
3  $$ -l h_rt=24:00:00
4  $$ -l killable
5  $$ -cwd
6  $$ -P KSU-GEN-RESERVED
7  $$ -q \*@@elves
8  $$ -pe mpi-spread 16
9
10 # Openmp 1
11 for i in 2 3 4
12 do
13 mpirun -np 16 /homes/mcheikh/CIS_625/hw5/part-1/Part1-MPI-OpenMP 1 1 1
14     $i $i
15 mpirun -np 8 /homes/mcheikh/CIS_625/hw5/part-1/Part1-MPI-OpenMP 1 1 1
16     $i $i
17 mpirun -np 4 /homes/mcheikh/CIS_625/hw5/part-1/Part1-MPI-OpenMP 1 1 1
18     $i $i
19 mpirun -np 2 /homes/mcheikh/CIS_625/hw5/part-1/Part1-MPI-OpenMP 1 1 1
20     $i $i
21
22 echo -e "-----DONE-----\n"
23 done

```

---

Listing 7: Part 1 - Bash script for OpenMP equal 2

---

```

1  #!/bin/bash
2  $$ -l mem=5G
3  $$ -l h_rt=24:00:00
4  $$ -l killable
5  $$ -cwd
6  $$ -P KSU-GEN-RESERVED
7  $$ -q \*@@elves
8  $$ -pe mpi-2 16
9
10 echo "PE_HOSTFILE:"
11 echo $PE_HOSTFILE
12 echo
13 echo "cat PE_HOSTFILE:"
14 cat $PE_HOSTFILE
15 echo '-----'
16
17 sed 's/ 2/ slots=1/g' $PE_HOSTFILE > NEW_HOSTFILE2
18 sed -i 's/gen-reserved.q@elf...beocat.ksu.edu UNDEFINED//g'
19     NEW_HOSTFILE2
20
21 # OpenMp = 2
22 for i in 2 3 4
23 do
24 mpirun --hostfile NEW_HOSTFILE2 -np 8 /homes/mcheikh/CIS_625/hw5/part
25     -1/Part1-MPI-OpenMP 2 1 1 $i $i
26 mpirun --hostfile NEW_HOSTFILE2 -np 4 /homes/mcheikh/CIS_625/hw5/part
27     -1/Part1-MPI-OpenMP 2 1 1 $i $i
28 mpirun --hostfile NEW_HOSTFILE2 -np 1 /homes/mcheikh/CIS_625/hw5/part
29     -1/Part1-MPI-OpenMP 2 1 1 $i $i

```

---

```

26
27 echo -e "-----DONE-----\n"
28 done

```

---

Listing 8: Part 1 - Bash script for OpenMP equal 4

---

```

1  #!/bin/bash
2  $$ -l mem=5G
3  $$ -l h_rt=24:00:00
4  $$ -l killable
5  $$ -cwd
6  $$ -P KSU-GEN-RESERVED
7  $$ -q \*@@elves
8  $$ -pe mpi-4 16
9
10 echo "PE_HOSTFILE:"
11 echo $PE_HOSTFILE
12 echo
13 echo "cat PE_HOSTFILE:"
14 cat $PE_HOSTFILE
15 echo '-----'
16
17 sed 's/ 4/ slots=1/g' $PE_HOSTFILE > NEW_HOSTFILE3
18 sed -i 's/gen-reserved.q@elf...beocat.ksu.edu UNDEFINED//g'
    NEW_HOSTFILE3
19
20 # OpenMp = 2
21 for i in 2 3 4
22 do
23 mpirun --hostfile NEW_HOSTFILE3 -np 4 /homes/mcheikh/CIS_625/hw5/part
    -1/Part1-MPI-OpenMP 4 1 1 $i $i
24 mpirun --hostfile NEW_HOSTFILE3 -np 1 /homes/mcheikh/CIS_625/hw5/part
    -1/Part1-MPI-OpenMP 4 1 1 $i $i
25
26 echo -e "-----DONE-----\n"
27 done

```

---